# Multi-View Unsupervised User Feature Embedding for Social Media-based Substance Use Prediction

**Tao Ding**
Department of Information Systems
University of Maryland, Baltimore County
taoding01@umbc.edu

**Warren K. Bickel**
Addiction Recovery Research Center
Virginia Tech Carilion Research Institute
wkbickel@vtc.vt.edu

**Shimei Pan**
Department of Information Systems
University of Maryland, Baltimore County
shimei@umbc.edu

## Abstract

In this paper, we demonstrate how the state-of-the-art machine learning and text mining techniques can be used to build effective social media-based substance use detection systems. Since a substance use ground truth is difficult to obtain on a large scale, to maximize system performance, we explore different unsupervised feature learning methods to take advantage of a large amount of unsupervised social media data. We also demonstrate the benefit of using multi-view unsupervised feature learning to combine heterogeneous user information such as Facebook "likes" and "status updates" to enhance system performance. Based on our evaluation, our best models achieved 86% AUC for predicting tobacco use, 81% for alcohol use and 84% for illicit drug use, all of which significantly outperformed existing methods. Our investigation has also uncovered interesting relations between a user's social media behavior (e.g., word usage) and substance use.

## 1 Introduction

A substance use disorder (SUD) is defined as a condition in which recurrent use of substances such as alcohol, drugs and tobacco causes clinically and functionally significant impairment in an individual's daily life (SAMHSA, 2015). According to the 2014 National Survey on Drug Use and Health, 1 in 10 Americans age 12 and older had a substance use disorder. Substance use also costs Americans more than $700 billion a year in increased health care costs, crimes and lost productivity (NIDA, 2015).

These days, people also spend a significant amount of time on social media such as Twitter, Facebook and Instagram to interact with friends and families, exchange ideas and thoughts, provide status updates and organize events and activities. The ubiquity and widespread use of social media underlines the needs to explore its intersection with substance use and its potential as a scalable and cost-effective solution for screening and preventing substance misuse and abuse.

In this research, we employ the state-of-the-art machine learning and text mining algorithms to build automated substance use prediction systems, which can be used to identify people who are at risk of SUD. Since SUD data are often expensive to obtain at a large scale, to maximize system performance, we focus on methods that employ unsupervised feature learning to take advantage of a large amount of unsupervised social media data. Previous research in Machine Learning, Image Processing, Speech and Natural language Processing has shown that to be able to utilize a large amount of unsupervised data is one of the most reliable ways to achieve good performance (Le et al., 2011; Lee et al., 2009; Le and Mikolov, 2014). Moreover, by analyzing rich human behavior data on social media, we can also gain insight into patterns of use and risk factors associated with substance use. The main contributions of this work include:

1. We have explored a comprehensive set of single-view feature learning methods to take advantage of a large amount of unsupervised social media data. Our results have shown significant improvement over baseline systems that only use supervised training data.

2. We have explored several multi-view learning algorithms to take advantage of heteroge-

neous user data such as Facebook "likes" and "status updates". Our results have demonstrated significant improvement over baselines that only use a single data type.

3. We have uncovered new insight into the relationship between a person's social media activities and substance use such as the relationship between word usage and SUD.

## 2 Related Work

Substance use disorder (SUD) encompasses a complex pattern of behaviors. Many studies have been conducted to discover factors interacting with SUD. A growing number of studies have confirmed a strong association between personal traits and substance use. For example, (Campbell et al., 2014) found that smokers have significantly higher *openness to experience* and lower *conscientiousness*, a personality trait related to a tendency to show self-discipline, act dutifully, and aim for achievement. (Cook et al., 1998) examined the links between alcohol consumption and personality and found that alcohol use is correlated positively with sociability and extraversion. (Terracciano et al., 2008) conducted a study involving 1102 participants and found a link between drug use and low *conscientiousness*. (Carroll et al., 2009) revealed risk factors related to addiction such as age, sex, impulsivity, sweet-liking, novelty reactivity, proclivity for exercise, and environmental impoverishment. Additionally, addiction is also linked to environmental and social factors such as neighborhood environment (Crum et al., 1996), family environment (Cadoret et al., 1986; Brent, 1995) and social norms (Botvin, 2000; Oetting and Beauvais, 1987).

Traditionally, in behavior science research, data are collected from surveys or interviews with a limited number of people. The advent of social media makes a large volume of diverse user data available to researchers, which makes it possible to study SUD based on online user behaviors in a natural setting. Typical data from social media include demographics (age, gender etc.), status updates (text posts etc.), social networks (follower and following graph etc.) and likes (thumb up/down etc.). Recently, social media analytics has increasingly become a powerful tool to help understand the traits and behaviors of millions of social media users such as personal traits (Gol-

beck et al., 2011; Volkova and Bachrach, 2015; Youyou et al., 2015; Kiliç and Pan, 2016), brand preferences (Yang et al., 2015), communities and events (Sayyadi et al., 2009), influenza trend (Aramaki et al., 2011) and crime (Li et al., 2012). So far, however, there has been limited work that directly applies large scale social media analytics to automatically predict SUD. Among the work known to us, (Zhou et al., 2016) identified common drug consumption behaviors with regard to the time of day and week. They also discovered common interests shared by drug users such as celebrities (e.g, Chris Tucker) and comedians (e.g., cheechandchong). In addition, (Kosinski et al., 2013) automatically predicted SUD based on social media likes. Since their dataset is very similar to ours, we will use the Kosinski model as one of our baselines.

## 3 Dataset

The data for the study was collected from 2007 to 2012 as a part of the myPersonality project (Kosinski et al., 2015). myPersonality was a popular Facebook application that offered to its users psychometric tests and feedback on their scores. The data were gathered with an explicit opt-in consent for reuse for research purposes. Our study uses three separate datasets from myPersonality: Facebook status updates (a.k.a. posts), Facebook likes and SUD status.

The *status update dataset* contains 22 million textual posts authored by 153,000 users. The average posts per user is 143 and the average words per user is 1730. We removed users who only have non-English posts and those who have written less than 500 words. Our final status update dataset includes 106,509 users with 21 million posts. After filtering out low frequency words (those appear less than 50 times in our corpus), the vocabulary size of the status update dataset is 73,935.

The *likes dataset* contains the Facebook likes used to express positive sentiment toward various targets such as products, movies, books, expressions, websites and people (they are called *Like Entities* or LEs). Previous studies have demonstrated that social media likes speak volumes about who we are. In addition to directly signaling interests and preferences, social media likes are indicative of ethnicity, intelligence and personality (Kosinski et al., 2013). The like dataset includes the likes of 11 million Facebook users.

Overall, there are 9.9 million unique LEs and 1.8 billion user-like pairs in this dataset. The average likes per user is 161 and the average Likes each LE received is 182. We filter out users who have a small number of likes as well as LEs receiving a small number of likes. The filtering threshold for users is 50 and is 800 for LEs. After the filtering, our like dataset contains 5,138,857 users and 253,980 unique LEs.

The SUD dataset contains a total of 13,557 participants (Stillwell and Tunney, 2012). Users were asked to answer questions like "Do you smoke?", with answers "daily or more", "less than daily" or "never". They also completed the Cigarette Dependence Scale (CDS-5) (Etter et al., 2003), Alcohol Use Questionnaire(AUQ) (Townshend and Duka, 2005) and the Assessment of Substance Misuse Questionnaire (ASMA) (Willner, 2000). Based on these assessments, the participants were divided into groups for each SUD type. For example, based on the assessment of tobacco use, a person is categorized as "daily or more" (group 3), "less than daily" (group 2), or "never" (group 1). The validity of the grouping was confirmed by the CDS-5 scores of the groups. Similarly, based on the assessment of alcohol use, participants were categorized as "weekly or more" (group 3), "less than once a week" (group 2) or "never" (group 1). Finally, based on the assessment of drug use, a person is assigned to "weekly or more" (group 3), "less than once a week" (group 2), or "never" (group 1). Among all the SUD participants, 37% of them are males and 63% are females. Their average age is 23 years old.

Since the like, status update and SUD datasets are only partially overlapping, their intersections are usually much smaller. Table 1 summarizes the sizes and usage of these datasets. Table 2 shows additional details of the SUD dataset including the distributions of each SUD class.

In summary, among all the datasets we have, the unsupervised like dataset is the largest (5 million+ people). We also have a significant amount of unsupervised status update data (100k+ users). In contrast, the supervised datasets which have the SUD ground truth are pretty small, ranging from 896 for the intersection of the likes, status updates and SUD (LikeStatusSUD in Table 1) to 3508, which is the intersection of the likes and SUD (LikesSUD in Table 1). Thus, the main focuses of this research include (1) employing unsupervised feature learning to take advantage of a large amount of unsupervised data (2) employing multi-view learning to combine heterogeneous user data for better prediction.

# 4 Single-View Post Embedding (SPE)

The main purpose of this study is to demonstrate the usefulness of employing unsupervised feature learning to derive a feature representation of a user's Facebook posts to take advantage of a large amount of unsupervised data. Since we only use Facebook status updates (a.k.a. posts) in this study, we call the process Single-view user Post Embedding (SPE).

## 4.1 SPE Feature Learning Methods

Since each user is associated with a sequence of textual posts, we have explored the following methods to learn a SPE for the user.

***Singular Value Decomposition (SVD)*** is a mathematical technique that is frequently used for dimension reduction (De Lathauwer et al., 2000). Given any $m * n$ matrix A, the algorithm will find matrices $U$, $V$ and $W$ such that $A = UWV^T$. Here $U$ is an orthonormal $m * n$ matrix, $W$ is a diagonal $n * n$ metrix and $V$ is an orthonormal $n * n$ matrix. Dimensionality reduction is done by computing $R = U * W_r$ where $W_r$ neglects all but the $r$ largest singular values in the diagonal matrix $W$. In our study, the $m$ is the number of users, $n$ is the number of unique words in the vocabulary. $A_{ij} = k$ where $k$ is how many times $word_j$ appears in $user_i$'s posts.

***Latent Dirichlet Allocation (LDA)*** is a generative graphical model that allows sets of documents to be explained by unobserved latent topics (Blei et al., 2003). For each document, LDA outputs a multinomial distribution over a set of latent topics. For each topic, LDA also outputs a multinomial distribution over the vocabulary.

To learn an SPE for each user based on all his/her posts, we have tried several methods (1) UserLDA: it treats all the posts from each user as one big document and trains an LDA model to drive the topic distribution for this document. The per-document topic distribution is then used as the SPE for this user. (2) PostLDA_Doc: it treats each post as a separate document and trains an LDA model to derive a topic distribution for each post. To derive the SPE for each user, we aggregate

Table 1: Dataset Descriptions

| Dataset | users | AvgUserLikes | AvgUserPosts | Usage |
|---|---|---|---|---|
| Likes | 5,138,857 | 184 | NA | Single View Feature Learning |
| LikesSUD | 3,508 | 267 | NA | Single View SUD Prediction |
| Status Update | 106,509 | NA | 143 | Single View Feature Learning |
| StatusSUD | 1,231 | NA | 195 | Single View SUD Prediction |
| LikeStatus | 54,757 | 232 | 220 | Multi-View Feature Learning |
| LikeStatusSUD | 896 | 277 | 219 | Multi-View SUD Predication |

Table 2: Class Distribution of Different SUD Datasets

| Dataset | Tabacco Use | | | Alcohol Use | | | Drug Use | | |
|---|---|---|---|---|---|---|---|---|---|
| | 3 | 2 | 1 | 3 | 2 | 1 | 3 | 2 | 1 |
| LikeSUD | 498 | 290 | 2603 | 469 | 1174 | 1716 | 171 | 276 | 1965 |
| StatusSUD | 226 | 95 | 880 | 179 | 416 | 596 | 76 | 102 | 671 |
| LikeStatusSUD | 147 | 69 | 660 | 123 | 290 | 453 | 262 | 53 | 75 |

all the per-post topic distribution vectors from the same user by averaging them. (3) PostLDA_Word: instead of using the $average$ of post-based topic distribution vectors, we used a word-based aggregation method suggested in (Schwartz et al., 2013):

$$p(topic|user) = \sum_{w \in voc} P(topic|w) * p(w|user)$$

where $voc$ represents the vocabulary, $p(w|user)$ is the probability that word $w$ appears in the posts of $user$ and $p(topic|w)$ is the topic distribution of a word $w$, which is available internally in an LDA model. For the UserLDA model, all the hyper parameters were set to default values. For all the PostLDA models, since Facebook posts are usually short and have a small number of topics in each post, we set the hyper parameter $\alpha$ to 0.3, as suggested in (Schwartz et al., 2013)

***Document Embedding with Distributed Memory (D-DM)*** Given a document, D-DM simultaneously learns a vector representation for each word and a vector for the entire document (Le and Mikolov, 2014). During training, the document vector and one or more word vectors are aggregated to predict a target word in the context. To learn an SPE for each user, we have explored two methods (1) User-D-DM: it treats all the posts by the same user as one document and trains a document vector to represent the user. (2) Post-D-DM: it treats each post as a document and train a D-DM to learn a vector for each post. To derive the SPE for a user, we aggregate all the post vectors from the same person using "average".

***Document Embedding with Distributed Bag of***

Table 3: SPE: Prediction Results

| Methods | Tobacco | Alcohol | Drug |
|---|---|---|---|
| Unigram | 0.663 | 0.672 | 0.644 |
| LIWC | 0.731 | 0.689 | 0.758 |
| SVD | 0.779 | 0.724 | 0.764 |
| UserLDA | 0.641 | 0.603 | 0.599 |
| PostLDA_Word | 0.733 | 0.617 | 0.628 |
| PostLDA_Doc | 0.768 | 0.687 | 0.721 |
| Post-D-DM | 0.536 | 0.622 | 0.520 |
| User-D-DM | 0.775 | 0.730 | 0.767 |
| Post-D-DBOW | 0.531 | 0.606 | 0.526 |
| User-D-DBOW | **0.802** | **0.768** | **0.819** |

***Words (D-DBOW)*** D-DBOW learns a global document vector to predict words randomly sampled from the document (Le and Mikolov, 2014). Unlike D-DM, D-DBOW only learns a vector for the entire document. It does not learn vectors for individual words. Neither does it use a local context window since the words for prediction are randomly sampled from the entire document. Similar to D-DM, to derive the SPE for a user, we used two methods (1) User-D-DBOW and (2) Post-D-DBOW.

## 4.2 SUD Prediction with SPE

In our experiments, to search for the best model, we systematically varied the output SPE dimension from 50, 100, 300, to 500. We used the Gensim implementation of SVD, LDA, D-DM and D-DBOW in our experiments. For D-DM, the context window size was set to 5.

We compared our models with two baselines that use only supervised learning (1) a unigram model which uses unigrams as the predicting features. Since we have a large number of unigrams, we performed supervised feature selection

to lower the total number of input features. Finally since all our SUD variables have three values, we employed SVM in 3-way classifications. (2) a LIWC model which uses human engineered LIWC features for SUD prediction. LIWC is a psycholinguistic lexicon (Pennebaker et al., 2015) that has been frequently used in text-based human behavior prediction. Since the number of LIWC features is relatively small, no feature selection was performed. Here, we only used the *Status Update* dataset in Table 1 as the training data for SPE learning and the *StatusSUD* dataset for supervised SUD prediction .

We evaluate the performance of our models using 10-fold cross validation. The evaluation results shown in Table 3 are based on weighted ROC AUC of the best models. Among all the feature learning methods for Facebook status updates, User-D-DBOW performed the best. It significantly outperformed all the baseline systems that only rely on supervised training ($p < 0.01$ based on t-tests). It also significantly outperformed all the traditional feature learning methods such as LDA and SVD ($p < 0.01$ based on t-tests). Moreover, in terms of whether to treat all the posts by the same user as one big document or separate documents, LDA prefers one post one document (models with a "post" prefix) while all the document vector-based methods prefer one user one document (models with a "User" prefix). Moreover, to use post-level LDA to derive the SPE of a user, the document-based aggregation method (PostLDA_Doc) performed better than the word-based method (PostLDA_Word).

## 5 Single-View Like Embedding (SLE)

In addition to textual posts, each user account is also associated with a set of likes. Since the like dataset is very sparse (e.g., among the millions of unique likes on Facebook, each user only has a small number of likes), we conduct experiments to learn a dense vector representation for all the likes by a user. We call this process Single-view user Like Embedding (SLE).

### 5.1 SLE Feature Learning Methods

The input to SLE is simply a set of LEs liked by a user. Each LE is represented by its id. To map such a representation to a dense user like vector, we have tried multiple methods.

*Singular Value Decomposition (SVD)* is similar to the one used in SPE except $A_{ij} = 1$ if $user_i$ likes $LE_j$. Otherwise, it is 0. Here $A$ is a $m * n$ matrix where $m$ is the number of users and $n$ is the number of unique LEs in the like dataset.

*Latent Dirichlet Allocation (LDA)*. To apply LDA to the like data, each individual LE is treated as a word token and all the LEs liked by the same person form a document. The order of the LEs in the document is random. For each user, LDA outputs a multinomial distribution over a set of latent "Like Topics". For example, a "Like Topic" about "hip hop music" may include famous hip hop songs and musicians.

*Autoencoder (AE)* is a neural network-based method for self-taught learning (Hinton and Salakhutdinov, 2006). It learns an identity function so that the output is as close to the input as possible. Although an identity function seems a trivial function to learn, by placing additional constraints (e.g,, to make the number of neurons in the hidden layer much smaller than that of the input), we can still force the system to uncover structures in the data. Architecturally, the AE we used has one input layer, one hidden layer and one output layer. For each user, we construct a training instance $(X, Y)$ where the input vector $X$ and output vector $Y$ are the same. The size of $X$ and $Y$ is the total number of unique LEs in our dataset. $X_i$ and $Y_i$ equal to 1 if the user likes $LE_i$. Otherwise they are 0.

*Document Vector with Distributed Memory (D-DM)* We also applied D-DM to the like data. Given all the likes of a user, D-DM learns a vector representation for each LE as well as a document vector for all the LEs from the same user. We use the learned document vector as the output SLE.

*Document Vector with Distributed Bag of Words (D-DBOW)* Similarly, we applied D-DBOW to the like dataset. Since D-DBOW does not use a local context window and the words for prediction are randomly sampled from the entire document, it is more appropriate for the like dataset than D-DM. where the relative positions of LEs do matter.

### 5.2 SUD Prediction with SLE

Similarly, we systematically varied the output SLE dimension from 50, 100, 300, to 500 in order to search for the best model. We used the Gensim implementation of SVD, LDA, D-DM and D-

Table 4: SLE: Prediction Results

| Method | Tobacco | Alcohol | Drug |
|---|---|---|---|
| Unigram | 0.687 | 0.651 | 0.673 |
| Kosinski* | 0.730* | 0.700* | 0.650* |
| AE | 0.678 | 0.648 | 0.672 |
| SVD | 0.757 | 0.756 | 0.753 |
| LDA | 0.723 | 0.737 | 0.704 |
| D-DM | 0.688 | 0.713 | 0.687 |
| D-DBOW | **0.787** | **0.795** | **0.791** |

⋆:2-way classification, 3-way for the others

DBOW in our experiments. For D-DM, the context window size was set to 20. We used Keras with Theano backend to implement AE.

We used SVM to perform 3-way classification. We compared our results with a unigram baseline. We also compared our results with that of the Kosinski model reported in (Kosinski et al., 2013). The Kosinski model was trained on the same Facebook like dataset. However, its results were based on two-way classification, a simpler task than 3-way classification. All the results are based on weighted ROC AUC.

As shown in Table 4, among all the SLE methods, the D-DBOW model performed the best. It significantly outperformed the unigram baseline that does not use any unsupervised data ($p < 0.01$ based on t-tests). It also significantly outperformed all the traditional feature learning method such as SVD and LDA (The Kosinski model used SVD for feature learning) ($p < 0.01$ based on t-tests). Between the two document vector-based methods, D-DBOW outperformed D-DM. We think this is due to the fact that D-DBOW does not use local context window, thus is not sensitive to the positions of LEs in a document. Since LE positions are randomly decided in our like data, D-DBOW seems to be a better fit for this dataset.

## 6 Multi-View User Embedding (MUE)

The main purpose of this study is to demonstrate the usefulness of combining heterogeneous user data such as likes and posts to learn a dense vector representation for each user. Since we employ unsupervised multi-view feature learning to combine these data, we call this process Multi-view User Embedding (MUE).

### 6.1 MUE Feature Learning Methods

We have explored two multi-view feature learning algorithms: Canonical Correlation Analysis (CCA) and Deep Canonical Correlation Analysis (DCCA).

***Canonical Correlation Analysis (CCA)*** CCA is a statistical method for exploring the relationships between two multivariate sets of variables (vectors) (Hardoon et al., 2004). Given two vectors $X$ and $Y$, CCA tries to find $aX$, $bY$ that are maximally correlated:

$$(a^*, b^*) = \arg\max_{a,b} corr(a^{'}X, b^{'}Y) \qquad (1)$$

$$= \arg\max_{a,b} \frac{a^{'}\sum_{XY} b}{\sqrt{a^{'}\sum_{XX} ab^{'}\sum_{YY} b}} \qquad (2)$$

where $(X, Y)$ denote random vectors with covariances $\sum_{XX} = Cov(X,X), \sum_{YY} = Cov(Y,Y)$ and cross-covariance $\sum_{XY} = Cov(X,Y)$. CCA has been used frequently in unsupervised data analysis (Sargin et al., 2006; Chaudhuri et al., 2009; Kumar and Daumé, 2011; Sharma et al., 2012).

***Deep Canonical Correlation Analysis (DCCA)*** DCCA aims to lean highly correlated deep architectures, which can be a non-linear extension of CCA (Andrew et al., 2013). The intuition is to find a maximally correlated representation of the two views by passing them through multiple stacked layers of nonlinear transformation (Andrew et al., 2013). Typically, there are three steps to train DCCA: (1) using a denoising autoencoder to pretrain each single view. In our experiments, we pretrain each single view using SPE or SLE. (2) computing the gradient of the correlation of top-level representation. (3) tuning parameters using back propagation to optimize the total correlation. Previously, DCCA was found to be more effective than CCA in image processing (Andrew et al., 2013).

### 6.2 SUD Prediction With MUE

The input to MUE are the two single views obtained earlier (i.e. SPE or SLE). Here, we choose the outputs from D-DBOW since it consistently outperformed all the other methods in learning SPEs and SLEs. We have run CCA and DCCA in two settings (1) balanced setting in which the SPE and SLE dimensions are always the same (2) imbalanced setting in which the dimension of SPE may be different from that of SLE. Since we varied the output dimensions of SPE and SLE from 50, 100, 300, to 500 systematically, the input dimension to MUE under the balanced setting are 100, 200, 600 and 1000. When running CCA

Figure 1: LIWC Features that are Most Significantly Correlated with Substance Use.

and DCCA under the imbalanced setting, we only chose the best SPE (with 50 dimensions) and the best SLE (with 300 dimensions). We also varied the number of MUE output dimensions systematically from 20,50,100,200,300,400,500 to 1000 (up to the total input MUE dimensions). We used the *LikeStatus* dataset in Table 1 as the training data for multi-view unsupervised feature learning. For MUE-based supervised SUD predication, we used the *LikeStatusSUD data*. In our experiments, we use a variant of CCA called *wGCCA* implemented by (Benton et al., 2016) where we set the weights for both views to be equal [1]. We used the DCCA implementation by (Andrew et al., 2013) which uses Keras and Theano as the deep learning platform [2]. We also varied the number of hidden layers from 1 to 3 to tune the performance.

We compared our multi-view learning results with 3 baselines: BestSPE and BestSLE are the best single view models. We also used a 3rd baseline called *Unigram_combine*, which simply concatenates all the post and like unigrams together and then applies supervised feature selection before uses the remaining features in a SVM-based classification. As shown in Table 5, both wGCCA and DCCA significantly outperformed the unigram-based baseline ($p < 0.01$ based on t-test). The difference between the best multi-view models (wGCCA_balanced for Alcohol and drug, wGCCA_imbalanced for illicit drugs) and the best single view models are also significant ($p < 0.01$). wGCCA also performed significantly better than DCCA on our tasks ($p < 0.01$ based on t-tests).

## 7 Social Media and Substance Use

In addition to building models that predict SUD, we are also interested in understanding the rela-

Table 5: MUE: Prediction Results

|  | Tobacco | Alcohol | Drug |
|---|---|---|---|
| BestSPE | 0.802 | 0.768 | 0.819 |
| BestSLE | 0.787 | 0.795 | 0.791 |
| Unigram_combine | 0.685 | 0.669 | 0.662 |
| wGCCA_balanced | 0.848 | **0.811** | **0.844** |
| wGCCA_imbalanced | **0.855** | 0.799 | 0.832 |
| DCCA_balanced | 0.774 | 0.778 | 0.742 |
| DCCA_imbalanced | 0.760 | 0.781 | 0.737 |

tionship between a person's social media activities and substance use behavior. Since many of the SPEs and SLEs are not easily interpretable, in this section, we focus on the LIWC features from status updates and the LDA topics from both Likes and status updates. Since the SUD ground truth is an ordinal variable and the LIWC/LDA features are numerical, we used Spearman's rank correlation analysis to identify features that are most significantly correlated with SUD. Figure 1 shows the LIWC features that are significantly correlated with at least one type of SUD ($p < 0.05$). The color red represents a positive correlation while blue represents a negative correlation. In addition, the saturation of the color indicates the significance of the correlation. The darker the color is, the more significant the correlation is.

As shown in Figure 1, swear words such as "fuck" and "shit", sexual words such as "horny" and "sex", words related to biological process such as "blood" and "pain" are positively correlated with all three types of SUD. In addition, words related to money such as "cash", words related to body such as "hands" and "legs", words related to ingestion such as "eat" and "drink" are positively correlated with both alcohol and drug use; words related to motion such as "car" and "go" are positively correlated with both alcohol and tobacco use. In addition, female references such as "girl" and "woman", prepositions, space reference words such as "up" and "down" are positively cor-

Table 6: Topics Most Significantly Correlated with Substance Use.

| | Significance | Topic |
|---|---|---|
| **Tobacco** | | |
| Posts | + | (T1) *fuck, shit, ass, fucking, bitch, face, don't, kick, damn, man, lol, hell...* |
| | - | (T2) *paper, book, writing, read, class, essay, english, finished, reading, time, page ...* |
| Likes | + | (T3) *Tool, Misfits, A Perfect Circle, Rob Zombie ...* |
| | - | (T4) *The Twilight Saga, Forever 21, Twilight, Victoria's Secret, Katy Perry* |
| **Alcohol** | | |
| Posts | + | (T5) *tonight, night, free, party, tickets, bar, saturday, friday, dj, drink, club, show, beer, ladies...* |
| | − | (T6) *class, history, paper, math, science, writing, essay, finished, study, test, final, exam ...* |
| Likes | ++ | (T7) *V For Vendetta, Boondock Saints, Pan's Labyrinth ...* |
| | − | (T8) *Cookie Monster, Squirt, Last Day of School, Hunger Games Official Page, Wonka ...* |
| **Drug** | | |
| Posts | ++ | (T9) *fuck, shit, ass, fucking, bitch, face, don't, kick, damn, man, lol, hell...* |
| | - | (T10) *dinner, nice, shopping, christmas, home, weekend, lunch, family, house,love,wine :-)...* |
| Likes | + | (T11) *Radiohead,The Cure, Depeche Mode, The Smiths, Arctic Monkeys ...* |
| | - | (T12) *Music, Movies, Traveling, Photography, Dancing ...* |

related with alcohol use, while words related to anger such as "hate" and "kill", words related to health such as "clinic" and "pill" are positively correlated with drug use.

In terms of LIWC features that are negatively correlated with SUD, words associated with the past such as "did" and "ago" are negatively correlated with both tobacco and drug use; assent words such as "ok", "yes" and "agree" are negatively correlated to both alcohol and tobacco use. In addition, male references such as "boy" and "man", words related to reward such as "prize" and "benefit", words related to positive emotions such as "nice" and "sweet", first person pronouns (plural) such as "we" and "our" are negatively correlated to drug use. Moreover, impersonal pronouns such as "it", differentiation words such as "but" and "else", and work-related words such as "job" and "work" are negatively correlated with alcohol use. Surprisingly, risk related words such as "danger", words related to sadness, death and negative emotions are also negatively correlated with alcohol use.

There are a few surprising correlations in our results. For example, female references such as "girl" and "woman" are positively related to alcohol use while male references such as "man" and "boy" are negatively related to drug use. To interpret this, previous research has shown (Schwartz et al., 2013) that female references actually are used more often by male authors and vice versa. Thus, our findings suggest that males are more likely to use alcohol while females are less likely

to use illicit drugs.

We have also used Spearman's correlation analysis to identify SUD-related "Like Topics" and "Status update Topics" learned by LDA. Since the number of significant topics is quite large, in Table 6, we only show a few samples. Based on a user's status updates, "swear topics" (T1, T9) are positively correlated with both tobacco and drug use, which is consistent with our LIWC findings. The "night life topic" (T5) is positively related to alcohol use. In addition, school related topics (T2, T6) are negatively correlated with tobacco and alcohol use. Positive family-related activities (T10) are negatively correlated with drug use. In addition, based on the LDA topics learned from "likes", a preference for rock music (T3,T11) is positive correlated with tobacco and drug use. A preference for movies such as "V For Vendetta" and "Boondock Saints" (T7) is positively correlated with alcohol use, while having a hobby (T12), liking cartoons and shows favored by kids (T8) or liking movies and brands favored by girls (T4) are negatively correlated with drug, alcohol and tobacco use.

## 8 Discussion and Future Work

Currently, our multi-view unsupervised features learning methods only learn from the intersection of the like and status update data, which is much smaller than either the like or the status update data. Similarly, MUE-based supervised prediction used only the intersection of all three datasets which is very small (only contains 896 users).

Thus, it would be useful if a future multi-view feature learning algorithm is capable of using all the available data (e.g., the union of all the supervised and unsupervised training data). Moreover, our best SPE model only has 50 dimensions while our best SLE model has 300 dimensions. This might be because the supervised training data used by SPE is almost three times smaller than that used by SLE . But surprisingly, SPE-based models performed better than SLE-based models. We expect that with more training data, the performance of SPE-based methods can be further improved.

## 9 Conclusion

We believe social media is a promising platform for both studying SUD-related human behaviors as well as engaging the public for substance abuse prevention and screening. In this study, we have focused on four main tasks (1) employing unsupervised features learning to take advantage of a large amount of unsupervised social media data (2) employing multi-view feature learning to combine heterogeneous user information such as "likes" and "status updates" to learn a comprehensive user representation (3) building SUD prediction models based on learned user features (4) employing correlation analysis to obtain human-interpretable results. Our investigation has not only produced models with the state-of-the-art prediction performance (e.g., for all three types of SUD, our models achieved over 80% prediction accuracy based on AUC), but also demonstrated the benefits of incorporating unsupervised heterogeneous user data for SUD prediction.

## References

Galen Andrew, Raman Arora, Jeff A Bilmes, and Karen Livescu. 2013. Deep canonical correlation analysis. In *ICML (3)*. pages 1247–1255.

Eiji Aramaki, Sachiko Maskawa, and Mizuki Morita. 2011. Twitter catches the flu: detecting influenza epidemics using Twitter. In *Proceedings of the conference on empirical methods in natural language processing*. Association for Computational Linguistics, pages 1568–1576.

Adrian Benton, Raman Arora, and Mark Dredze. 2016. Learning multiview embeddings of Twitter users. In *Proceedings of the 54th Annual Meeting of the Association for Computational Linguistics*. volume 2, pages 14–19.

David M Blei, Andrew Y Ng, and Michael I Jordan. 2003. Latent dirichlet allocation. *Journal of machine Learning research* 3(Jan):993–1022.

Gilbert J Botvin. 2000. Preventing drug abuse in schools: Social and competence enhancement approaches targeting individual-level etiologic factors. *Addictive behaviors* 25(6):887–897.

David A Brent. 1995. Risk factors for adolescent suicide and suicidal behavior: mental and substance abuse disorders, family environmental factors, and life stress. *Suicide and Life-Threatening Behavior* 25(s1):52–63.

Remi J Cadoret, Ed Troughton, Thomas W O'Gorman, and Ellen Heywood. 1986. An adoption study of genetic and environmental factors in drug abuse. *Archives of general psychiatry* 43(12):1131–1136.

S Campbell, L Henry, J Hammelman, and M Pignatore. 2014. Personality and smoking behaviour of non-smokers, previous smokers, and habitual smokers. *J Addict Research & Therapy* 5:191.

Marilyn E Carroll, Justin J Anker, and Jennifer L Perry. 2009. Modeling risk factors for nicotine and other drug abuse in the preclinical laboratory. *Drug and alcohol dependence* 104:S70–S78.

Kamalika Chaudhuri, Sham M Kakade, Karen Livescu, and Karthik Sridharan. 2009. Multi-view clustering via canonical correlation analysis. In *Proceedings of the 26th annual international conference on machine learning*. ACM, pages 129–136.

Mark Cook, Alison Young, Dean Taylor, and Anthony P Bedford. 1998. Personality correlates of alcohol consumption. *Personality and Individual Differences* 24(5):641–647.

Rosa M Crum, Marsha Lillie-Blanton, and James C Anthony. 1996. Neighborhood environment and opportunity to use cocaine and other drugs in late childhood and early adolescence. *Drug and alcohol dependence* 43(3):155–161.

Lieven De Lathauwer, Bart De Moor, and Joos Vandewalle. 2000. A multilinear singular value decomposition. *SIAM journal on Matrix Analysis and Applications* 21(4):1253–1278.

Jean-Francois Etter, Jacques Le Houezec, and Thomas V Perneger. 2003. A self-administered questionnaire to measure dependence on cigarettes: the cigarette dependence scale. *Neuropsychopharmacology* 28(2):359.

Jennifer Golbeck, Cristina Robles, and Karen Turner. 2011. Predicting personality with social media. In *CHI'11 extended abstracts on human factors in computing systems*. ACM, pages 253–262.

David R Hardoon, Sandor Szedmak, and John Shawe-Taylor. 2004. Canonical correlation analysis: An overview with application to learning methods. *Neural computation* 16(12):2639–2664.

G. Hinton and R. Salakhutdinov. 2006. Reducing the dimensionality of data with neural networks. *science* 313(5786):504–507.

Işil Doğa Yakut Kiliç and Shimei Pan. 2016. Analyzing and preventing bias in text-based personal trait prediction algorithms. In *Tools with Artificial Intelligence (ICTAI), 2016 IEEE 28th International Conference on*. IEEE, pages 1060–1067.

Michal Kosinski, Sandra C Matz, Samuel D Gosling, Vesselin Popov, and David Stillwell. 2015. Facebook as a research tool for the social sciences: Opportunities, challenges, ethical considerations, and practical guidelines. *American Psychologist* 70(6):543.

Michal Kosinski, David Stillwell, and Thore Graepel. 2013. Private traits and attributes are predictable from digital records of human behavior. *Proceedings of the National Academy of Sciences* 110(15):5802–5805.

Abhishek Kumar and Hal Daumé. 2011. A co-training approach for multi-view spectral clustering. In *Proceedings of the 28th International Conference on Machine Learning (ICML-11)*. pages 393–400.

Quoc V Le and Tomas Mikolov. 2014. Distributed representations of sentences and documents. In *ICML*. volume 14, pages 1188–1196.

Quoc V Le, Will Y Zou, Serena Y Yeung, and Andrew Y Ng. 2011. Learning hierarchical invariant spatio-temporal features for action recognition with independent subspace analysis. In *Computer Vision and Pattern Recognition (CVPR), 2011 IEEE Conference on*. IEEE, pages 3361–3368.

Honglak Lee, Peter Pham, Yan Largman, and Andrew Y Ng. 2009. Unsupervised feature learning for audio classification using convolutional deep belief networks. In *Advances in neural information processing systems*. pages 1096–1104.

Rui Li, Kin Hou Lei, Ravi Khadiwala, and Kevin Chen-Chuan Chang. 2012. Tedas: A Twitter-based event detection and analysis system. In *2012 IEEE 28th International Conference on Data Engineering*. IEEE, pages 1273–1276.

NIDA. 2015. Drugs, brains, and behavior: The science of addiction. https://www.drugabuse.gov/publications/drugs-brains-behavior-science-addiction/introduction.

ER Oetting and Fred Beauvais. 1987. Common elements in youth drug abuse: Peer clusters and other psychosocial factors. *Journal of Drug Issues* 17(2):133–151.

James W Pennebaker, Ryan L Boyd, Kayla Jordan, and Kate Blackburn. 2015. The development and psychometric properties of liwc2015. Technical report.

SAMHSA. 2015. Substance use disorders. https://www.samhsa.gov/disorders/substance-use.

Mehmet Emre Sargin, Engin Erzin, Yücel Yemez, and A Murat Tekalp. 2006. Multimodal speaker identification using canonical correlation analysis. In *Acoustics, Speech and Signal Processing, 2006. ICASSP 2006 Proceedings. 2006 IEEE International Conference on*. IEEE, volume 1, pages I–I.

Hassan Sayyadi, Matthew Hurst, and Alexey Maykov. 2009. Event detection and tracking in social streams. In *ICWSM*.

H Andrew Schwartz, Johannes C Eichstaedt, Margaret L Kern, Lukasz Dziurzynski, Stephanie M Ramones, Megha Agrawal, Achal Shah, Michal Kosinski, David Stillwell, Martin EP Seligman, et al. 2013. Personality, gender, and age in the language of social media: The open-vocabulary approach. *PloS one* 8(9):e73791.

Abhishek Sharma, Abhishek Kumar, Hal Daume, and David W Jacobs. 2012. Generalized multiview analysis: A discriminative latent space. In *Computer Vision and Pattern Recognition (CVPR), 2012 IEEE Conference on*. IEEE, pages 2160–2167.

David J Stillwell and Richard J Tunney. 2012. Effects of measurement methods on the relationship between smoking and delay reward discounting. *Addiction* 107(5):1003–1012.

Antonio Terracciano, Corinna E Löckenhoff, Rosa M Crum, O Joseph Bienvenu, and Paul T Costa. 2008. Five-factor model personality profiles of drug users. *Bmc Psychiatry* 8(1):22.

Julia M Townshend and Theodora Duka. 2005. Binge drinking, cognitive performance and mood in a population of young social drinkers. *Alcoholism: Clinical and Experimental Research* 29(3):317–325.

Svitlana Volkova and Yoram Bachrach. 2015. On predicting sociodemographic traits and emotions from communications in social networks and their implications to online self-disclosure. *Cyberpsychology, Behavior, and Social Networking* 18(12):726–736.

Paul Willner. 2000. Further validation and development of a screening instrument for the assessment of substance misuse in adolescents. *Addiction* 95(11):1691–1698.

Chao Yang, Shimei Pan, Jalal Mahmud, Huahai Yang, and Padmini Srinivasan. 2015. Using personal traits for brand preference prediction. In *EMNLP*. pages 86–96.

Wu Youyou, Michal Kosinski, and David Stillwell. 2015. Computer-based personality judgments are more accurate than those made by humans. *Proceedings of the National Academy of Sciences* 112(4):1036–1040.

Yiheng Zhou, Numair Sani, Chia-Kuei Lee, and Jiebo Luo. 2016. Understanding illicit drug use behaviors by mining social media. *arXiv preprint arXiv:1604.07096* .